

Advanced Machine Learning

Assignment 1: Neural Networks

Osama Bin Zahir

```
In [113... %%capture
# Installing required packages
!pip install tensorflow
!pip install tensorflow-datasets
```

```
In [114... (imdb_a_train, imdb_b_train), (imdb_a_test, imdb_b_test) = imdb.load_data(num_words=10000)
```

```
In [115... # Importing required Libraries

import numpy as np
from tensorflow.keras.datasets import imdb
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

```
In [116... max([max(sequence) for sequence in imdb_a_train])
```

```
Out[116]: 9999
```

```
In [117... # Preparing the data for the model

# Retrieving a dictionary mapping words to their index in the IMDB dataset
word_index = imdb.get_word_index()

# Creating a reverse dictionary and mapping the indices back to the original words
reverse_word_index = dict(
    [(value, key) for (key, value) in word_index.items()])
# Create models with different configurations
# Decoding reviews from the IMDB dataset
decoded_review = " ".join([reverse_word_index.get(i - 3, "?") for i in imdb_a_train[0]
```

```
In [118... # Encoding the integer sequences via multi-hot encoding
```

```
import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
imdb_train = vectorize_sequences(imdb_a_train)
imdb_test = vectorize_sequences(imdb_a_test)
```

```
In [119... imdb_train[0]
```


Epoch 1/20
30/30 [=====] - 3s 43ms/step - loss: 0.1534 - accuracy: 0.8093 - val_loss: 0.1060 - val_accuracy: 0.8742

Epoch 2/20
30/30 [=====] - 0s 14ms/step - loss: 0.0769 - accuracy: 0.9163 - val_loss: 0.0891 - val_accuracy: 0.8898

Epoch 3/20
30/30 [=====] - 0s 13ms/step - loss: 0.0560 - accuracy: 0.9423 - val_loss: 0.0854 - val_accuracy: 0.8893

Epoch 4/20
30/30 [=====] - 0s 13ms/step - loss: 0.0436 - accuracy: 0.9601 - val_loss: 0.0841 - val_accuracy: 0.8869

Epoch 5/20
30/30 [=====] - 0s 13ms/step - loss: 0.0349 - accuracy: 0.9707 - val_loss: 0.0855 - val_accuracy: 0.8848

Epoch 6/20
30/30 [=====] - 0s 13ms/step - loss: 0.0283 - accuracy: 0.9782 - val_loss: 0.0863 - val_accuracy: 0.8811

Epoch 7/20
30/30 [=====] - 0s 13ms/step - loss: 0.0236 - accuracy: 0.9833 - val_loss: 0.0876 - val_accuracy: 0.8795

Epoch 8/20
30/30 [=====] - 0s 13ms/step - loss: 0.0199 - accuracy: 0.9877 - val_loss: 0.0896 - val_accuracy: 0.8790

Epoch 9/20
30/30 [=====] - 0s 13ms/step - loss: 0.0168 - accuracy: 0.9899 - val_loss: 0.0915 - val_accuracy: 0.8756

Epoch 10/20
30/30 [=====] - 0s 13ms/step - loss: 0.0140 - accuracy: 0.9924 - val_loss: 0.0928 - val_accuracy: 0.8764

Epoch 11/20
30/30 [=====] - 0s 14ms/step - loss: 0.0120 - accuracy: 0.9938 - val_loss: 0.0943 - val_accuracy: 0.8737

Epoch 12/20
30/30 [=====] - 0s 17ms/step - loss: 0.0103 - accuracy: 0.9945 - val_loss: 0.0960 - val_accuracy: 0.8721

Epoch 13/20
30/30 [=====] - 1s 17ms/step - loss: 0.0089 - accuracy: 0.9953 - val_loss: 0.0973 - val_accuracy: 0.8716

Epoch 14/20
30/30 [=====] - 1s 18ms/step - loss: 0.0079 - accuracy: 0.9959 - val_loss: 0.0988 - val_accuracy: 0.8705

Epoch 15/20
30/30 [=====] - 1s 18ms/step - loss: 0.0071 - accuracy: 0.9963 - val_loss: 0.1004 - val_accuracy: 0.8680

Epoch 16/20
30/30 [=====] - 1s 17ms/step - loss: 0.0064 - accuracy: 0.9965 - val_loss: 0.1007 - val_accuracy: 0.8689

Epoch 17/20
30/30 [=====] - 1s 17ms/step - loss: 0.0058 - accuracy: 0.9967 - val_loss: 0.1018 - val_accuracy: 0.8678

Epoch 18/20
30/30 [=====] - 1s 18ms/step - loss: 0.0053 - accuracy: 0.9969 - val_loss: 0.1028 - val_accuracy: 0.8665

Epoch 19/20
30/30 [=====] - 1s 19ms/step - loss: 0.0049 - accuracy: 0.9970 - val_loss: 0.1039 - val_accuracy: 0.8654


Epoch 20/20
30/30 [=====] - 1s 17ms/step - loss: 0.0046 - accuracy: 0.9971 - val_loss: 0.1044 - val_accuracy: 0.8660

```
In [125... history_dict = history_one_hidden_layer.history
history_dict.keys()
```

```
Out[125]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [126... # Plotting the training and validation loss


import matplotlib.pyplot as plt
history_dict = history_one_hidden_layer.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
In [127... # Plotting the training and validation accuracy

plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
```

```
plt.legend()
plt.show()
```



In [128...

```
results = model_one_hidden_layer.evaluate(imdb_test,value_test)
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.1132 - accuracy: 0.8554
```

In [129...

```
# Adding Dropout Layer & Regularizers
```

```
from tensorflow import keras
from tensorflow.keras import layers
from keras.layers import Dense
from tensorflow.keras import regularizers
from keras.layers import Dropout
```

```
model_one_hidden_layer = Sequential()
model_one_hidden_layer.add(Dense(32, activation='tanh', activity_regularizer=regularizers.l2(1e-4)))
model_one_hidden_layer.add(Dropout(0.5))
model_one_hidden_layer.add(Dense(1, activation='sigmoid'))
```

```
# compiling the model
```

```
model_one_hidden_layer.compile(optimizer="adam",
                               loss="mean_squared_error",
                               metrics=["accuracy"])
```

```
# Validating the model
```

```
a_val = imdb_train[:10000]
```

[illegible]

Epoch 1/20
30/30 [=====] - 2s 47ms/step - loss: 0.1853 - accuracy: 0.77
97 - val_loss: 0.1419 - val_accuracy: 0.8633

Epoch 2/20
30/30 [=====] - 0s 15ms/step - loss: 0.1202 - accuracy: 0.90
12 - val_loss: 0.1214 - val_accuracy: 0.8882

Epoch 3/20
30/30 [=====] - 0s 17ms/step - loss: 0.0999 - accuracy: 0.93
16 - val_loss: 0.1153 - val_accuracy: 0.8889

Epoch 4/20
30/30 [=====] - 1s 18ms/step - loss: 0.0866 - accuracy: 0.94
89 - val_loss: 0.1128 - val_accuracy: 0.8869

Epoch 5/20
30/30 [=====] - 1s 17ms/step - loss: 0.0772 - accuracy: 0.96
01 - val_loss: 0.1124 - val_accuracy: 0.8825

Epoch 6/20
30/30 [=====] - 1s 17ms/step - loss: 0.0699 - accuracy: 0.96
83 - val_loss: 0.1137 - val_accuracy: 0.8772

Epoch 7/20
30/30 [=====] - 1s 19ms/step - loss: 0.0644 - accuracy: 0.97
52 - val_loss: 0.1156 - val_accuracy: 0.8739

Epoch 8/20
30/30 [=====] - 1s 19ms/step - loss: 0.0596 - accuracy: 0.97
87 - val_loss: 0.1163 - val_accuracy: 0.8737

Epoch 9/20
30/30 [=====] - 1s 23ms/step - loss: 0.0553 - accuracy: 0.98
31 - val_loss: 0.1187 - val_accuracy: 0.8682

Epoch 10/20
30/30 [=====] - 1s 19ms/step - loss: 0.0515 - accuracy: 0.98
67 - val_loss: 0.1202 - val_accuracy: 0.8657

Epoch 11/20
30/30 [=====] - 1s 20ms/step - loss: 0.0485 - accuracy: 0.98
88 - val_loss: 0.1224 - val_accuracy: 0.8622

Epoch 12/20
30/30 [=====] - 1s 21ms/step - loss: 0.0460 - accuracy: 0.99
01 - val_loss: 0.1251 - val_accuracy: 0.8561

Epoch 13/20
30/30 [=====] - 1s 28ms/step - loss: 0.0437 - accuracy: 0.99
16 - val_loss: 0.1276 - val_accuracy: 0.8521

Epoch 14/20
30/30 [=====] - 1s 26ms/step - loss: 0.0410 - accuracy: 0.99
35 - val_loss: 0.1294 - val_accuracy: 0.8499

Epoch 15/20
30/30 [=====] - 1s 30ms/step - loss: 0.0393 - accuracy: 0.99
49 - val_loss: 0.1322 - val_accuracy: 0.8473

Epoch 16/20
30/30 [=====] - 1s 29ms/step - loss: 0.0375 - accuracy: 0.99
55 - val_loss: 0.1342 - val_accuracy: 0.8447

Epoch 17/20
30/30 [=====] - 1s 23ms/step - loss: 0.0360 - accuracy: 0.99
60 - val_loss: 0.1366 - val_accuracy: 0.8399


Epoch 18/20
30/30 [=====] - 1s 26ms/step - loss: 0.0350 - accuracy: 0.99
67 - val_loss: 0.1389 - val_accuracy: 0.8369

Epoch 19/20
30/30 [=====] - 1s 23ms/step - loss: 0.0335 - accuracy: 0.99
67 - val_loss: 0.1420 - val_accuracy: 0.8324

Epoch 20/20
30/30 [=====] - 1s 22ms/step - loss: 0.0328 - accuracy: 0.99
67 - val_loss: 0.1425 - val_accuracy: 0.8329

In [130...

```
plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



In [131...

```
results = model_one_hidden_layer.evaluate(imdb_test,value_test)
```

782/782 [=====] - 3s 3ms/step - loss: 0.1526 - accuracy: 0.8188

In [132...

```
# Adding more hidden layers to examine how it affects the accuracy. Here three hidden

from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras import regularizers

model_three_hidden_layers = Sequential()
model_three_hidden_layers.add(Dense(32, activation='tanh', activity_regularizer=regularizers.l2(1e-4)))
model_three_hidden_layers.add(Dropout(0.5))
model_three_hidden_layers.add(Dense(32, activation='tanh', activity_regularizer=regularizers.l2(1e-4)))
model_three_hidden_layers.add(Dropout(0.5))
model_three_hidden_layers.add(Dense(32, activation='tanh', activity_regularizer=regularizers.l2(1e-4)))
model_three_hidden_layers.add(Dropout(0.5))
model_three_hidden_layers.add(Dense(1, activation='sigmoid')) # Output layer
```


1

[illegible]

Epoch 1/20
30/30 [=====] - 5s 85ms/step - loss: 0.2446 - accuracy: 0.70
59 - val_loss: 0.1605 - val_accuracy: 0.8634

Epoch 2/20
30/30 [=====] - 1s 23ms/step - loss: 0.1744 - accuracy: 0.87
83 - val_loss: 0.1401 - val_accuracy: 0.8788

Epoch 3/20
30/30 [=====] - 1s 31ms/step - loss: 0.1466 - accuracy: 0.90
87 - val_loss: 0.1371 - val_accuracy: 0.8754

Epoch 4/20
30/30 [=====] - 1s 27ms/step - loss: 0.1294 - accuracy: 0.92
55 - val_loss: 0.1372 - val_accuracy: 0.8711

Epoch 5/20
30/30 [=====] - 1s 32ms/step - loss: 0.1162 - accuracy: 0.93
97 - val_loss: 0.1398 - val_accuracy: 0.8646

Epoch 6/20
30/30 [=====] - 1s 27ms/step - loss: 0.1084 - accuracy: 0.94
57 - val_loss: 0.1395 - val_accuracy: 0.8621

Epoch 7/20
30/30 [=====] - 1s 29ms/step - loss: 0.1013 - accuracy: 0.95
06 - val_loss: 0.1398 - val_accuracy: 0.8581

Epoch 8/20
30/30 [=====] - 1s 24ms/step - loss: 0.0947 - accuracy: 0.95
65 - val_loss: 0.1409 - val_accuracy: 0.8594

Epoch 9/20
30/30 [=====] - 1s 30ms/step - loss: 0.0892 - accuracy: 0.96
07 - val_loss: 0.1420 - val_accuracy: 0.8531

Epoch 10/20
30/30 [=====] - 1s 24ms/step - loss: 0.0851 - accuracy: 0.96
33 - val_loss: 0.1426 - val_accuracy: 0.8509

Epoch 11/20
30/30 [=====] - 1s 25ms/step - loss: 0.0806 - accuracy: 0.96
65 - val_loss: 0.1464 - val_accuracy: 0.8454

Epoch 12/20
30/30 [=====] - 1s 27ms/step - loss: 0.0747 - accuracy: 0.97
29 - val_loss: 0.1483 - val_accuracy: 0.8422

Epoch 13/20
30/30 [=====] - 1s 31ms/step - loss: 0.0721 - accuracy: 0.97
37 - val_loss: 0.1467 - val_accuracy: 0.8440

Epoch 14/20
30/30 [=====] - 1s 23ms/step - loss: 0.0687 - accuracy: 0.97
66 - val_loss: 0.1482 - val_accuracy: 0.8407

Epoch 15/20
30/30 [=====] - 1s 28ms/step - loss: 0.0655 - accuracy: 0.97
72 - val_loss: 0.1498 - val_accuracy: 0.8381

Epoch 16/20
30/30 [=====] - 1s 24ms/step - loss: 0.0618 - accuracy: 0.97
99 - val_loss: 0.1503 - val_accuracy: 0.8377

Epoch 17/20
30/30 [=====] - 1s 27ms/step - loss: 0.0602 - accuracy: 0.98
15 - val_loss: 0.1518 - val_accuracy: 0.8332


Epoch 18/20
30/30 [=====] - 1s 30ms/step - loss: 0.0586 - accuracy: 0.98
28 - val_loss: 0.1530 - val_accuracy: 0.8332

Epoch 19/20
30/30 [=====] - 1s 39ms/step - loss: 0.0564 - accuracy: 0.98
31 - val_loss: 0.1546 - val_accuracy: 0.8298

Epoch 20/20
30/30 [=====] - 1s 43ms/step - loss: 0.0541 - accuracy: 0.98
58 - val_loss: 0.1539 - val_accuracy: 0.8304

In [133...

```
plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



In [134...

```
results = model_three_hidden_layers.evaluate(imdb_test,value_test)
```

782/782 [=====] - 2s 3ms/step - loss: 0.1644 - accuracy: 0.8114

In [135...

```
# Building a model with 64 hidden units and one hidden layer to examine the accuracy

from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras import regularizers

model_one_hidden_layer_64_units = Sequential()
model_one_hidden_layer_64_units.add(Dense(64, activation='tanh', activity_regularizer=
model_one_hidden_layer_64_units.add(Dropout(0.5))
model_one_hidden_layer_64_units.add(Dense(1, activation='sigmoid')) # Output Layer

# Compiling the model

model_one_hidden_layer_64_units.compile(optimizer="adam",
loss="mean_squared_error",
```

[illegible]

Epoch 1/20
30/30 [=====] - 3s 53ms/step - loss: 0.1677 - accuracy: 0.8076 - val_loss: 0.1254 - val_accuracy: 0.8801

Epoch 2/20
30/30 [=====] - 1s 25ms/step - loss: 0.1033 - accuracy: 0.9212 - val_loss: 0.1160 - val_accuracy: 0.8900

Epoch 3/20
30/30 [=====] - 1s 35ms/step - loss: 0.0849 - accuracy: 0.9448 - val_loss: 0.1147 - val_accuracy: 0.8839

Epoch 4/20
30/30 [=====] - 1s 38ms/step - loss: 0.0722 - accuracy: 0.9629 - val_loss: 0.1154 - val_accuracy: 0.8790

Epoch 5/20
30/30 [=====] - 1s 35ms/step - loss: 0.0639 - accuracy: 0.9725 - val_loss: 0.1179 - val_accuracy: 0.8742

Epoch 6/20
30/30 [=====] - 1s 44ms/step - loss: 0.0576 - accuracy: 0.9781 - val_loss: 0.1206 - val_accuracy: 0.8689

Epoch 7/20
30/30 [=====] - 1s 45ms/step - loss: 0.0520 - accuracy: 0.9839 - val_loss: 0.1238 - val_accuracy: 0.8636

Epoch 8/20
30/30 [=====] - 1s 39ms/step - loss: 0.0480 - accuracy: 0.9879 - val_loss: 0.1268 - val_accuracy: 0.8584

Epoch 9/20
30/30 [=====] - 1s 34ms/step - loss: 0.0448 - accuracy: 0.9889 - val_loss: 0.1294 - val_accuracy: 0.8553

Epoch 10/20
30/30 [=====] - 1s 38ms/step - loss: 0.0422 - accuracy: 0.9915 - val_loss: 0.1327 - val_accuracy: 0.8488

Epoch 11/20
30/30 [=====] - 1s 33ms/step - loss: 0.0394 - accuracy: 0.9931 - val_loss: 0.1383 - val_accuracy: 0.8404

Epoch 12/20
30/30 [=====] - 1s 38ms/step - loss: 0.0374 - accuracy: 0.9943 - val_loss: 0.1413 - val_accuracy: 0.8349

Epoch 13/20
30/30 [=====] - 1s 36ms/step - loss: 0.0360 - accuracy: 0.9941 - val_loss: 0.1431 - val_accuracy: 0.8359

Epoch 14/20
30/30 [=====] - 1s 37ms/step - loss: 0.0344 - accuracy: 0.9952 - val_loss: 0.1465 - val_accuracy: 0.8326

Epoch 15/20
30/30 [=====] - 1s 34ms/step - loss: 0.0329 - accuracy: 0.9965 - val_loss: 0.1482 - val_accuracy: 0.8293

Epoch 16/20
30/30 [=====] - 1s 36ms/step - loss: 0.0312 - accuracy: 0.9971 - val_loss: 0.1515 - val_accuracy: 0.8268

Epoch 17/20
30/30 [=====] - 1s 35ms/step - loss: 0.0302 - accuracy: 0.9971 - val_loss: 0.1552 - val_accuracy: 0.8240


Epoch 18/20
30/30 [=====] - 1s 38ms/step - loss: 0.0294 - accuracy: 0.9967 - val_loss: 0.1570 - val_accuracy: 0.8221

Epoch 19/20
30/30 [=====] - 1s 37ms/step - loss: 0.0284 - accuracy: 0.9973 - val_loss: 0.1595 - val_accuracy: 0.8186

Epoch 20/20
30/30 [=====] - 1s 43ms/step - loss: 0.0278 - accuracy: 0.9974 - val_loss: 0.1610 - val_accuracy: 0.8176

In [136...

```
plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



In [137...

```
results = model_one_hidden_layer_64_units.evaluate(imdb_test,value_test)
```

782/782 [=====] - 3s 3ms/step - loss: 0.1711 - accuracy: 0.8002

- Neural Network Configurations for IMDB Sentiment Analysis
- Introduction

This report presents the results of training several neural network models with different configurations for sentiment analysis on the IMDB dataset. The goal is to determine the optimal model architecture for accurately classifying movie reviews as positive or negative.

- Experimental Setup

Dataset: The IMDB dataset consists of 50,000 movie reviews labeled as positive or negative.

Model Configurations: Four different neural network configurations were evaluated: One hidden

layer with 32 units and Tanh activation function One hidden layer with 32 units, Tanh activation function, and dropout regularization Three hidden layers with 32 units each, Tanh activation function, and dropout regularization One hidden layer with 64 units, Tanh activation function, and dropout regularization

- Results and Analysis
- Model 1: One Hidden Layer (32 Units, Tanh Activation) Accuracy: 85.54% Validation Accuracy: 86.60% Observations: The model shows a good performance with decent accuracy on both training and validation sets. However, there is a slight overfitting as the training accuracy is slightly higher than the validation accuracy.
- Model 2: One Hidden Layer (32 Units, Tanh Activation, Dropout Regularization) Accuracy: 81.88% Validation Accuracy: 83.29% Observations: Introducing dropout regularization slightly reduces overfitting compared to Model 1, but there is still room for improvement in validation accuracy.
- Model 3: Three Hidden Layers (32 Units, Tanh Activation, Dropout Regularization) Accuracy: 81.14% Validation Accuracy: 83.04% Observations: Adding more hidden layers does not significantly improve performance. The model seems to suffer from overfitting, similar to Model 2.
- Model 4: One Hidden Layer (64 Units, Tanh Activation, Dropout Regularization) Accuracy: 80.02% Validation Accuracy: 81.76% Observations: Increasing the number of units in the hidden layer does not lead to better performance. The model exhibits similar overfitting issues as previous configurations.

Model	Validation Accuracy	Test Loss	Test Accuracy
One Hidden Layer (32 units)	0.8660	0.1132	0.8554
One Hidden Layer (32 units)	0.8329	0.1526	0.8188
Three Hidden Layers (32 units)	0.8304	0.1644	0.8114
One Hidden Layer (64 units)	0.8176	0.1711	0.8002

- Conclusion

Based on the experiments conducted, the model with one hidden layer consisting of 32 units and Tanh activation function (Model 1) performs the best, achieving the highest accuracy on both the training and validation sets. Introducing dropout regularization helps mitigate overfitting to some extent, but adding more hidden layers or increasing the number of units does not yield significant improvements. Therefore, Model 1 is recommended as the optimal configuration for sentiment analysis on the IMDB dataset. However, further experimentation with hyperparameter tuning and different architectures could potentially lead to even better results.